

Lab 5 - Linux Memory Managment

1 – Setup & Basic Tasks

Check the number of CPUs and the number of online-cpus (using which command?)

- nproc
- lscpu
- cat /proc/cpuinfo

Check the compiler installation (using which command?)

- gcc --version

Setup three to four terminal connections to your machine, it makes life easier

Tip: use **tmux**

Look up for a cheat sheet, but here are a few useful commands:

Combination	Description
tmux	Create a tmux session
CTRL+b into c	Create a terminal
CTRL+b into %	Split current pane vertically
CTRL+b into "	Split current pane horizontally
CTRL+b into <nr>	Switch to existing terminal <nr>
CTRL+d into <nr>	Close terminal <nr>

1.1 – Basic memory under Linux

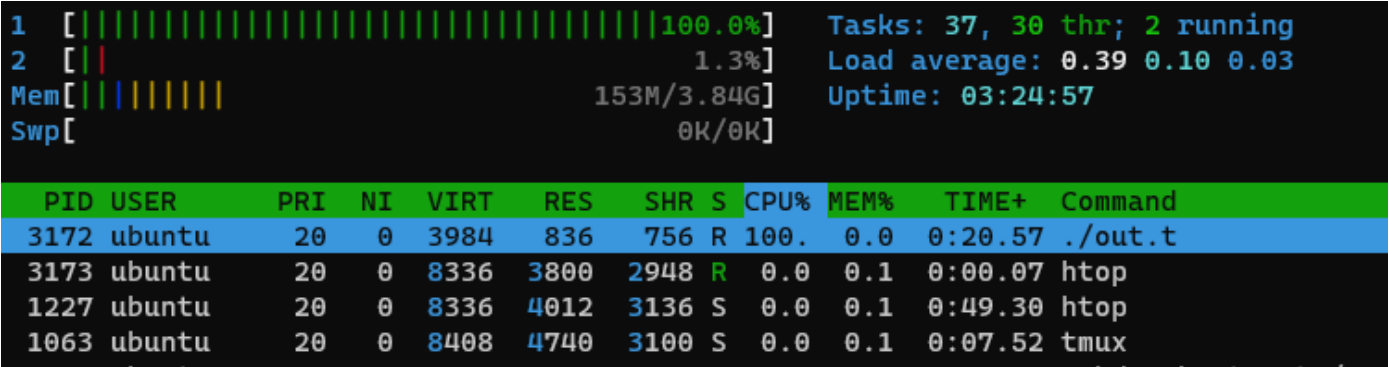
Todo 1 - In the first section call up a function to print out the PID of the running process - this will come in useful later (hint - use the getpid system call). Add an

endless loop below this, exit, compile and run.

See [MEM_lab.c](#) - Todo 1

In a second terminal run top or htop - what memory parameters are useful to know? (Hint - use the documentation for htop to look for VIRT/RES/SHR)

```
$ htop
```



Tip: Press M to order htop by memory usage.

Column	Description
M_SIZE (VIRT)	The size of the virtual memory of the process.
M_RESIDENT (RES)	The resident set size (text + data + stack) of the process (i.e. the size of the process's used physical memory).
M_SHARE (SHR)	The size of the process's shared pages.

In a third terminal using the command free (hint - man free) display the system memory parameters in kilobytes

```
$ free --kilo
```

	total	used	free	shared	buff/cache	available
Mem:	4127260	162123	3318194	991	646942	3718672
Swap:	0	0	0			

Tip: Use e.g. -k to display kibibytes (1024 or 2^10 bytes) instead of --kilo for kilobytes (1000 or 10^3 bytes).

Read the file `/proc/${pid}/status` specifically the memory related portions. What do the fields mean? (hint-man proc)

Attribute	Description
Mems_allowed	Mask of memory nodes allowed to this process.
RssAnon	Size of resident anonymous memory.
RssShmem	Size of resident shared memory.
VmLck	Locked memory size.
VmPeak	Peak virtual memory size.
VmPin	These are pages that can't be moved because something needs to directly access physical memory.
VmSize	Virtual memory size.
VmSwap	Swapped-out virtual memory size.

VmLck

The amount of memory that has been "pinned" or kept in physical memory and is not allowed to be swapped out to disk. This is often used for critical system processes or processes that require low-latency access to memory.

VmPin

The amount of memory that has been locked or "locked-in" physical memory, meaning it cannot be paged out or swapped out to disk. This is typically used for processes that need to guarantee access to a certain amount of memory, such as databases or real-time applications.

What is the difference between VmPin and VmLck?

Attribute	Description
VmLck	Locked memory size.
VmPin	These are pages that can't be moved because something needs to directly access physical memory.

Research the command smem, install if necessary. What information does smem give you about your system?

smem is used to report memory usage with shared memory divided proportionally.

PID	User	Command	Swap	USS	PSS	RSS
3172	ubuntu	./out.t	0	208	260	1872
3148	ubuntu	tmux	0	596	1063	3440
3173	ubuntu	htop	0	852	1166	3800
1227	ubuntu	htop	0	880	1210	4016
3174	ubuntu	-bash	0	1720	1933	5104
1085	ubuntu	-bash	0	1720	1937	5100
3159	ubuntu	-bash	0	1720	1937	5076
1105	ubuntu	-bash	0	1724	1947	5092
3149	ubuntu	-bash	0	1764	1990	5308
1095	ubuntu	-bash	0	1776	2009	5212
3110	ubuntu	-bash	0	1800	2029	5324
1064	ubuntu	-bash	0	1848	2084	5456
1063	ubuntu	tmux	0	1864	2324	4744
901	ubuntu	/lib/systemd/systemd --user	0	1532	2838	9472
3185	ubuntu	/usr/bin/python3 /usr/bin/s	0	6244	8179	12620

Column	Description
USS	Unique Set Size for unshared memory
PSS	Proportional Set Size = shared + unshared memory
RSS	Resident Set Size: Held in RAM, will not swap to disk

Todo 2 - Remove the endless loop and insert code to read the page size (store it in a variable) and print it out.

See [MEM_lab.c](#) - Todo 2 and output at the end of Task 1.1.

Verify this with the getconf command

```
$ getconf PAGE_SIZE
4096
```

Todo 3 - Include code to reserve memory (hint - man malloc) the size of a number of pages.

See [MEM_lab.c](#) - Todo 3 and output at the end of Task 1.1.

After each execution step of this code run the following command (another terminal)

```
ps -o min_flt,maj_flt {pid}
```

See PS output at the end of Task 1.1.

Attribute	Description
min_flt	Number of minor page faults
max_flt	Number of major page faults

Todo 4 - use the `align_alloc` function to reserve a buffer of a number of pages size, aligned on a page boundary. Then use the function `mincore` to check whether the pages are in memory.

The pages are not resident in the memory before accessing them.

Here we write random characters into the memory as displayed with the samples.

See [MEM_lab.c](#) - Todo 4 and output at the end of Task 1.1.

Todo 5 - Linux uses lazy allocation - include an access to the buffer - and run the code again. What do you see when you run the code and check the page faults reported by the `ps` command?

The count for page faults reset when starting the program anew as a new process starts and different memory is allocated.

See [MEM_lab.c](#) - Todo 5 (included in Part 4) and output at the end of Task 1.1.

OUTPUT: Terminal of complete program

```

Hello MEM Lab
----- Part 1: Simple program check memory of process
----- the PID of this process is 4700

Press enter to continue

----- Part 2: whats the page size
Page size: 4096
Press return to continue

----- Part 3: generate some memory area
Allocated 16384 bytes of memory at address 0x55c54984fac0

Press return to continue

----- Part 4: Are these pages in memory?
Allocated 16384 bytes of memory at address 0x55c549854000
mincore:
Before buffer access:
Page 0: not in memory
Page 1: not in memory
Page 2: not in memory
Page 3: not in memory
Memory samples:
After buffer access:
Page 0: in memory
Page 1: in memory
Page 2: in memory
Page 3: in memory
Memory samples: abcdefghijklmnopqrstuvwxyzabcdefghijklmnopqrstuvwxyzabcdefghijkl
Press return to continue

----- Part 5: Lets limit the available memory
Check the number of page faults
Press return to continue

Carry out the instructions in the lab guide to limit the available memory and repeat
Bye MEM Lab

```

OUTPUT: PS of complete program

```

ubuntu@ubuntu:~/08/P05$ ps -o minflt,majflt 4700
MINFLT  MAJFL
      89      0
ubuntu@ubuntu:~/08/P05$ ps -o minflt,majflt 4700
MINFLT  MAJFL
      89      0
ubuntu@ubuntu:~/08/P05$ ps -o minflt,majflt 4700
MINFLT  MAJFL
      90      0
ubuntu@ubuntu:~/08/P05$ ps -o minflt,majflt 4700
MINFLT  MAJFL
      96      0
ubuntu@ubuntu:~/08/P05$ ps -o minflt,majflt 4700
MINFLT  MAJFL
    37637      0

```

1.2 – Limiting memory (1)

From reading the process status file we know the maximum amount of memory the process uses during startup. We can now attempt to limit this on a high level.

Research the `ulimit` command and use it to display the resource limitations.

What precisely is limited?

The following command `ulimit -a` fetches all limits for a user.

```
elia@elia-ThinkPad-X1-Nano-Gen-1:~/文档/Ausbildung/Schule/Studium/ZHAW/6. Semest  
er/BSY/Lab/P05$ ulimit -a  
real-time non-blocking time (microseconds, -R) unlimited  
core file size              (blocks, -c) 0  
data seg size               (kbytes, -d) unlimited  
scheduling priority         (-e) 0  
file size                   (blocks, -f) unlimited  
pending signals             (-i) 62372  
max locked memory           (kbytes, -l) 2012356  
max memory size             (kbytes, -m) unlimited  
open files                  (-n) 1024  
pipe size                   (512 bytes, -p) 8  
POSIX message queues        (bytes, -q) 819200  
real-time priority          (-r) 0  
stack size                  (kbytes, -s) 8192  
cpu time                    (seconds, -t) unlimited  
max user processes          (-u) 62372  
virtual memory              (kbytes, -v) unlimited  
file locks                  (-x) unlimited
```

Using the data from reading `/proc/${pid}/status` let us limit the available memory for the start phase of the test program to under the peak requirement. What happens?

For example. If we limit a process that requires 90 MB (92'160 kB) and we limit it to 50 MB (51'200 kb) using the command `ulimit -m 51200`, the program will receive an error when trying to allocate more than 50 MB and therefore will crash.

How do you restore the unlimited memory access capability? What is your assessment of this method?

- `ulimit -m unlimited`
- It might be helpful in the first time to assign unlimited memory access. But if a process has flaws, it might be dangerous.

1.3 – Limiting memory (2)

cgroups allows us to limit the resources used for individual processes. Here we will create a memory controller for our test process

We define a cgroup memory controller. Use the man pages to understand the parameters

```
sudo cgcreate -a ubuntu:ubuntu -t ubuntu:ubuntu -g  
memory:myGroup
```

Parameter	Description
sudo	Run command with admin privileges
cgcreate	Create a cgroup
-a ubuntu:ubuntu	Define user & group to access the cgroup as admins. User: Ubuntu, Group: Ubuntu
-t ubuntu:ubuntu	Define user & group that the cgroup will apply to. User: Ubuntu, Group: Ubuntu
-t memory:myGroup	Define subsystem and cgroup name. Subsystem: Memory, Name: myGroup

What memory controller files have been created? What can be used to set limitations of memory usage?

- The following files have been created
 - `memory.limit_in_bytes`
 - `memory.usage_in_bytes`
 - `memory.max_usage_in_bytes`
 - `memory.failcnt`
- What can be used to set limitations of memory usage?

- For example we could use:

- `sudo cgset -r memory.limit_in_bytes=500000 myGroup`

We check the peak memory usage of the process which should be, in section 5, high - around the 120M mark.

We can now use this value to limit the process memory by writing an appropriate value into the group memory controller file

```
echo 2M > /sys/fs/cgroup/memory/myGroup/memory.limit_in_bytes
```

By running the process as follows ...

```
cgexec -g memory:myGroup process_name
```

... the process will run under the condition set by the cgroups memory controller.

Check the results using the `free` command. Two things can happen - if it's your lucky day the process will be killed. Why? The console output of your VM will give you a better hint. Explain it.

- The process will be killed because it has exceeded the memory limit set by myGroup. OOM comes into action and kills the process.
- Use the command `dmesg` to display the system log:

```
[615982.516840] Memory cgroup out of memory: Killed process 16823 (out.t) total-vm:19852kB, anon-rss:1848kB, file-rss:1812kB, shmem-rss:0kB, UID:1000 pgtables:52kB oom_score_adj:0
```

1.4 – Setting up a swap area

Why should we bother with a swap area? Because the principle of virtual memory depends on having excess secondary memory to enable a maximum number of processes to run “simultaneously.”

If using `free` it can be seen there is no swap area in secondary memory, then one needs to be setup. We do this in the following sequence

- 1.) Create a file that can be used for swapping
 - a.) `sudo fallocate -l 1G /swapfile`
- 2.) Give this file root permissions only

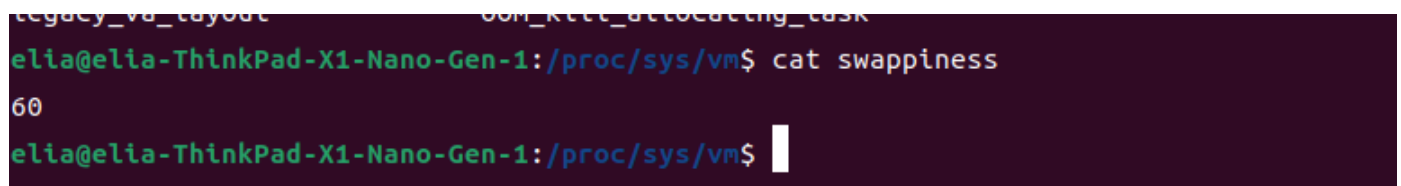
- a.) `sudo chmod 600 /swapfile`
- 3.) Setup a Linux swap area in the file
 - a.) `sudo mkswap /swapfile`
- 4.) Activate the swap file
 - a.) `sudo swapon /swapfile`
- 5.) If this is to be permanent then
 - a.) `sudo nano /etc/fstab`
 - b.) And add: `/swapfile swap swap defaults 0 0`
- 6.) `sudo swapon --show` or 7.) `sudo free -h` will now show a swap area

Swappiness is a Linux kernel property that defines how often the system will use the swap space.

Use the command to read the swappiness.

```
cat /proc/sys/vm/swappiness
```

What value do you get? The higher the value the more likely the kernel is to swap, the lower the value the more the kernel tries to avoid it.



```
elia@elia-ThinkPad-X1-Nano-Gen-1:/proc/sys/vm$ cat swappiness
60
elia@elia-ThinkPad-X1-Nano-Gen-1:/proc/sys/vm$
```

On production servers, a low swappiness is often preferred to decrease latencies and user available system memory.

Swappiness can be adjusted using:

```
sudo sysctl vm.swappiness=10
```

Read the manpage for `swapon`

If your process was killed in Substep 1.3, setup the swap area and repeat the experiment. What happens now?

The swaps swapped a lot due to swapping the swappiness configuration of the memory for more swappiness. Swap :)

	total	used	free	shared	buff/cache	available
Mem:	4030528	196008	2416532	972	1417988	3589212
Swap:	1048572	76800	971772			